

# Using Computer Clusters with Python

BIOS 274: Introductory Python Programming for Genomics

Veronica Behrens

12/9/2019

# What is a computer cluster?

A group of powerful computers (and other resources like software and storage) that act like a single system

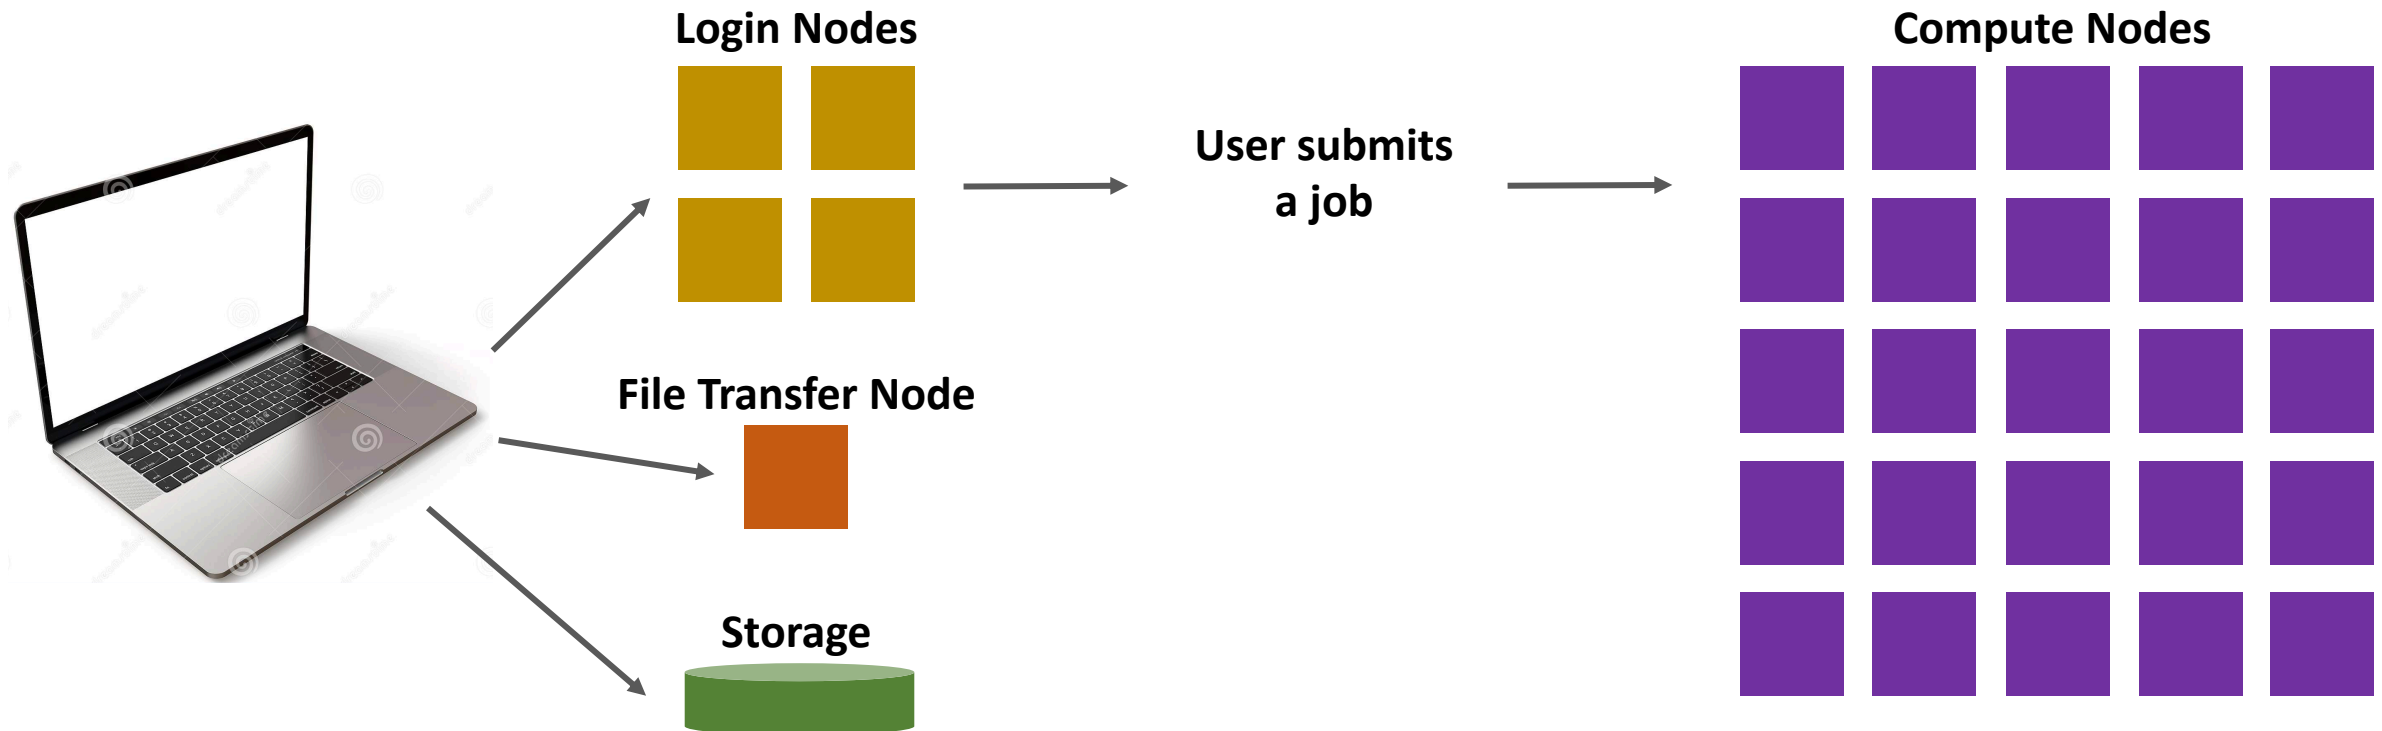


# Why use a computer cluster?

- For many genomics “big data” analyses, your local laptop/computer is not powerful enough!
- On the cluster:
  - Useful software is already installed
  - There’s lots of storage
  - There’s lots of RAM
  - You can parallelize jobs

# How are computer clusters organized?

- **Login/Head nodes:** For copying data, editing scripts, short test runs
- **Compute nodes:** For running jobs (scripts)
- **Storage**



# Computer clusters at Stanford

**Available to everyone at Stanford for free:**

- FarmShare (rice)

**Must be added as an authorized user by PI:**

- SCG4
- Sherlock

# Logging in to the FarmShare cluster

**ssh SUNetID@rice.stanford.edu**

- password
- DUO authentication

# Transferring files to and from the cluster

- **Secure File Transfer Protocol**

- `sftp`
- `lpwd`, `lcd`, `lls` to navigate local computer (your laptop)
- `pwd`, `cd`, `ls` to navigate the remote computer (the cluster / rice)
- `put FILENAME` to transfer to cluster from local computer
- `get FILENAME` to transfer from cluster to local computer

- **Secure Copy**

- `scp FILE_ON_LAPTOP SUNetID@rice.stanford.edu:DIRECTORY_ON_CLUSTER`
- `scp SUNetID@rice.stanford.edu:FILE_ON_CLUSTER DIRECTORY_ON_LAPTOP`

- **User-Friendly Graphical User Interfaces (GUIs)**

- Fetch (Mac): <https://fetchsoftworks.com/fetch/>
- CyberFX (Windows): [https://uit.stanford.edu/software/scrt\\_sfx](https://uit.stanford.edu/software/scrt_sfx)

# Accessing software on the cluster

<b>module avail</b>	List all available modules
<b>module load MODULE_NAME</b>	Load a module for use
<b>module unload MODULE_NAME</b>	Unload a module
<b>module list</b>	List all currently loaded modules
<b>module spider MODULE_NAME</b>	Detailed information about a particular module

**Check out all the software in `/usr/bin`, too!**



# Some useful software

- **bedtools**
- **samtools**
- bamtools
- vcftools
- **bcftools**
- sratoolkit
- picard-tools     manipulating various sequencing data files
- ucsc\_tools     manipulating various sequencing data files
- **blastn**     BLAST on nucleotides
- muscle     multiple sequence alignment
- **mafft**     multiple sequence alignment

**bold** indicates the tool is available on rice

# Some useful software

- Dealing with FASTQ files and mapping
  - bwa mapping (DNA-seq)
  - **bowtie2** mapping (DNA-seq)
  - **STAR** mapping (RNA-seq)
  - salmon mapping (RNA-seq)
  - **fastqc** stats about fastq file
  - trim\_galore quality filter fastq files
  - **trimmomatic** quality filter fastq files
  - cutadapt quality filter fastq files

**bold** indicates the tool is available on rice

# Submitting a job on rice using the SLURM job scheduler

## 1. Save the following text in a file called `testJob.sh`

```
#!/bin/bash
#SBATCH --job-name=testJob
#SBATCH --mail-type=END,FAIL
#SBATCH --mail-user=SUNetID@stanford.edu
#SBATCH --ntasks=1
#SBATCH --mem-per-cpu=1GB
#SBATCH --time=00:05:00
#SBATCH --output=testJob.log

##### YOUR COMMANDS TO RUN HERE #####
echo 'Running test job!' > testJob.txt
python3 test.py
```

# Job name  
# Mail events (NONE, BEGIN, END, FAIL, ALL)  
# Where to send notification emails  
# Number of CPUs to use  
# Job memory request  
# Time limit in hrs:min:sec  
# Output file for job

## 2. Submit the script as job: `sbatch testJob.sh`

### Other useful SLURM commands:

- `squeue -u SUNetID` Check the status of your jobs
- `scancel JOB_ID` Delete a job

# Login to rice, make directory and symbolic links

1. `ssh SUNetID@rice.stanford.edu`
2. `mkdir Day9`
3. `cd Day9`
4. `ln -s /afs/ir.stanford.edu/class/gene211/misc/BIOS274/Day9/blastQuery.fa`
5. `ln -s /afs/ir.stanford.edu/class/gene211/misc/BIOS274/Day9/hg38.chrom.sizes`
6. `ln -s /afs/ir.stanford.edu/class/gene211/misc/BIOS274/Day9/hg38_exons.bed`
7. `ln -s /afs/ir.stanford.edu/class/gene211/misc/BIOS274/Day9/hg38_genes.bed`
8. `ln -s /afs/ir.stanford.edu/class/gene211/misc/BIOS274/Day9/hg38.fa.gz`
9. `ln -s /afs/ir.stanford.edu/class/gene211/misc/BIOS274/Day9/test.bam`

# Run blast from the command line

# Make a blast database from the hg38 genome

# After running this command, we will also have hg38.nhr, hg38.nin, and hg38.nsq files

```
zcat hg38.fa.gz | makeblastdb -dbtype nucl -out hg38 -title hg38
```

# Run blast

```
blastn -db hg38 -query blastQuery.fa -out blastQuery_results.txt
```

# How can we make the output more useful?

```
blastn -db hg38 -query blastQuery.fa -out blastQuery_results.tsv \  
-outfmt '6 sseqid sstart send qseqid pident'
```

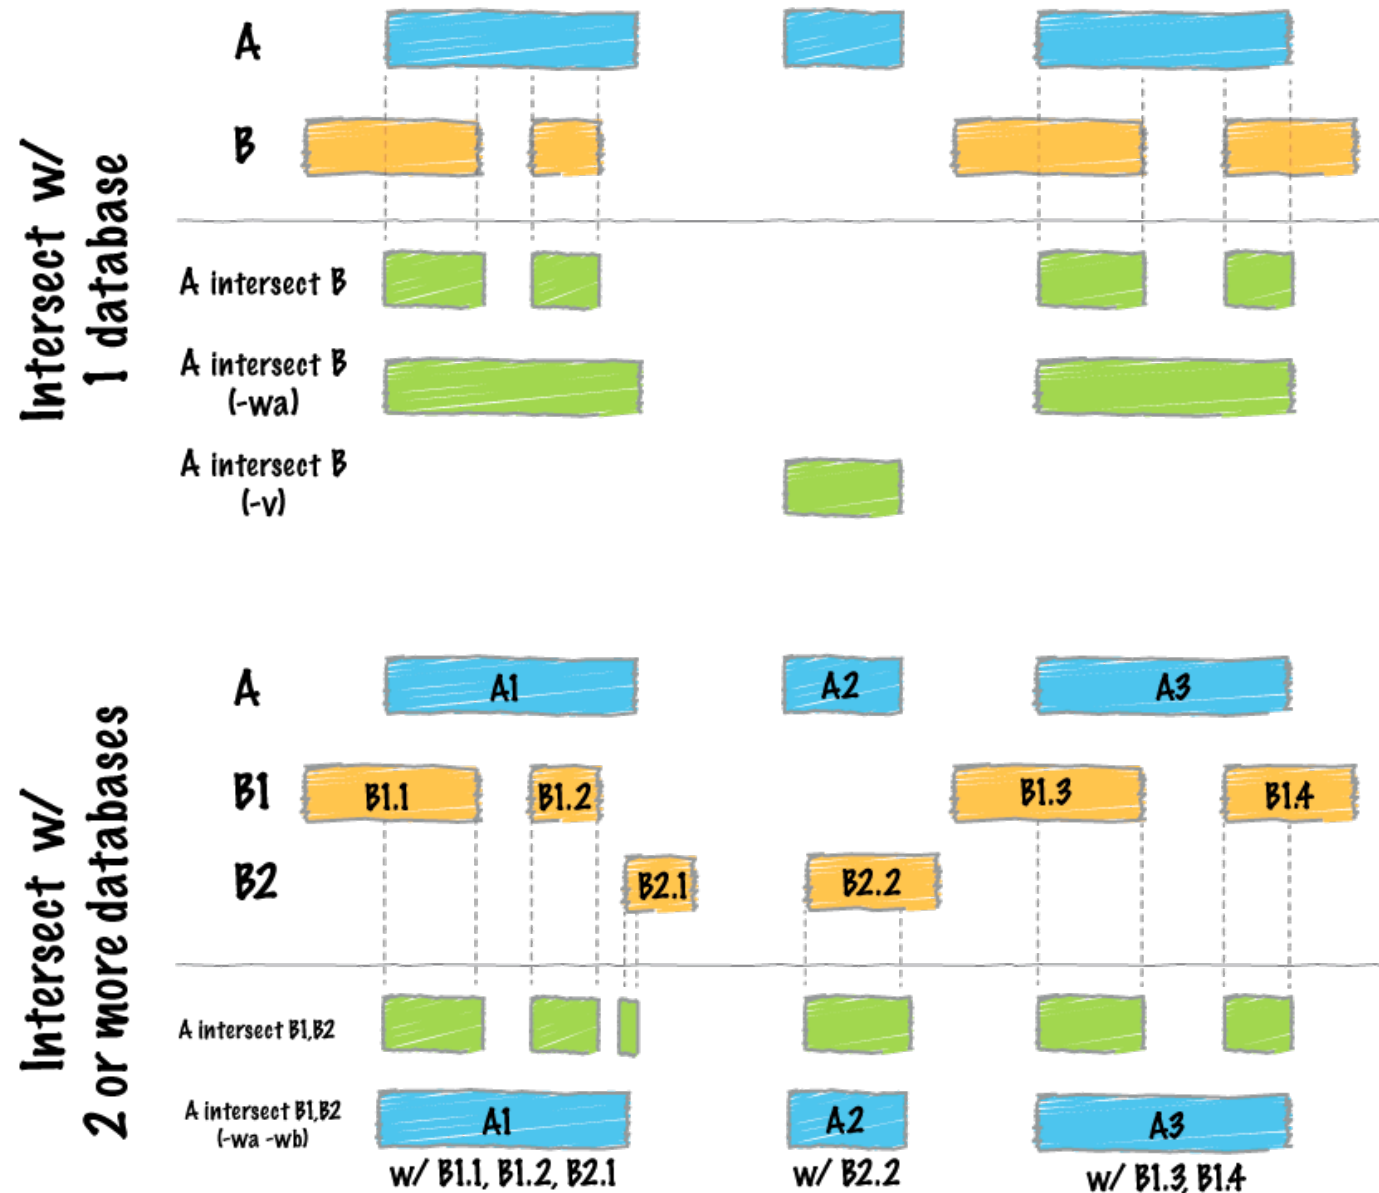
# Write a python script called formatResults.py that removes 'chr' from the beginning of each chromosome name, ensures the start <= end, and condenses the final 2 columns into a single |-delimited column

```
python3 formatResults.py blastQuery_results.tsv > blastQuery_results.bed
```

# Which queries are in exons?

```
bedtools intersect -wa -a blastQuery_results.bed -b hg38_exons.bed \  
> blastQuery_results_inExons.bed
```

# bedtools intersect



# Download ChIP-seq data

## # Download: ZFP91 ChIP-seq on human K562

```
wget https://www.encodeproject.org/files/ENCFF150ZBH/@@download/ENCFF150ZBH.bed.gz
gunzip ENCFF150ZBH.bed.gz
mv ENCFF150ZBH.bed ZFP91_K562.bed
```

## # Download: EGR1 ChIP-seq on human K562

```
wget https://www.encodeproject.org/files/ENCFF175VSS/@@download/ENCFF175VSS.bed.gz
gunzip ENCFF175VSS.bed.gz
mv ENCFF175VSS.bed EGR1_K562.bed
```

## # Download: FOXA1 ChIP-seq on human K562

```
wget https://www.encodeproject.org/files/ENCFF765NAN/@@download/ENCFF765NAN.bed.gz
gunzip ENCFF765NAN.bed.gz
mv ENCFF765NAN.bed FOXA1_K562.bed
```

## # Download: K562 DNase-seq

```
wget https://www.encodeproject.org/files/ENCFF821KDJ/@@download/ENCFF821KDJ.bed.gz
gunzip ENCFF821KDJ.bed.gz
mv ENCFF821KDJ.bed DNase_K562.bed
```

# Find the most interesting regions

# Find regions that overlap an EGR1, a FOXA1, a ZFP91, and a DNase peak.

```
bedtools intersect -a EGR1_K562.bed -b FOXA1_K562.bed | \
bedtools intersect -a stdin -b ZFP91_K562.bed | \
bedtools intersect -a stdin -b DNase_K562.bed > \
EGR1_FOXA1_ZFP91_DNase_K562.bed
```

# Make the format more useful.

```
python3 formatResults.py EGR1_FOXA1_ZFP91_DNase_K562.bed \
> EGR1_FOXA1_ZFP91_DNase_K562_temp.bed
```

```
mv EGR1_FOXA1_ZFP91_DNase_K562_temp.bed EGR1_FOXA1_ZFP91_DNase_K562.bed
```

# Which gene is closest to each of these regions?

```
bedtools sort -faidx hg38.chrom.sizes -i EGR1_FOXA1_ZFP91_DNase_K562.bed \
> EGR1_FOXA1_ZFP91_DNase_K562_sorted.bed
```

```
mv EGR1_FOXA1_ZFP91_DNase_K562_sorted.bed EGR1_FOXA1_ZFP91_DNase_K562.bed \
bedtools closest -a EGR1_FOXA1_ZFP91_DNase_K562.bed -b hg38_genes.bed \
-g hg38.chrom.sizes
```

# Use cut to generate a file with only the most important columns.

```
cut -f1,2,3,8,9 EGR1_FOXA1_ZFP91_DNase_K562_withGenes.bed \
> EGR1_FOXA1_ZFP91_DNase_K562_withGenes_simple.bed
```



# Let's work with a BAM file!

# Try to look at test.bam. What happens? Can samtools help?

```
samtools view test.bam
```

# Let's look at some informative stats about test.bam

```
samtools flagstat test.bam
```

# How many reads have a MAPQ score above 30?

```
samtools view -q 30 test.bam | wc -l
```

# What is the read depth in these three regions?

```
chr1:10311-10472770
```

```
chr1:16212-16230346
```

```
chr1:104944-104970612
```

```
samtools index test.sam
```

```
samtools bedcov roi.bed test.sam
```